

# Deploying a Microservices Application on Amazon EKS (Elastic Kubernetes Service)

## Table of Contents

1. Problem Statements
2. AWS Services Used
3. Architecture Diagram
4. Implementation Steps
5. Testing & Validation
6. Clean Up AWS Resources
7. Conclusion

## 1. Problem Statement

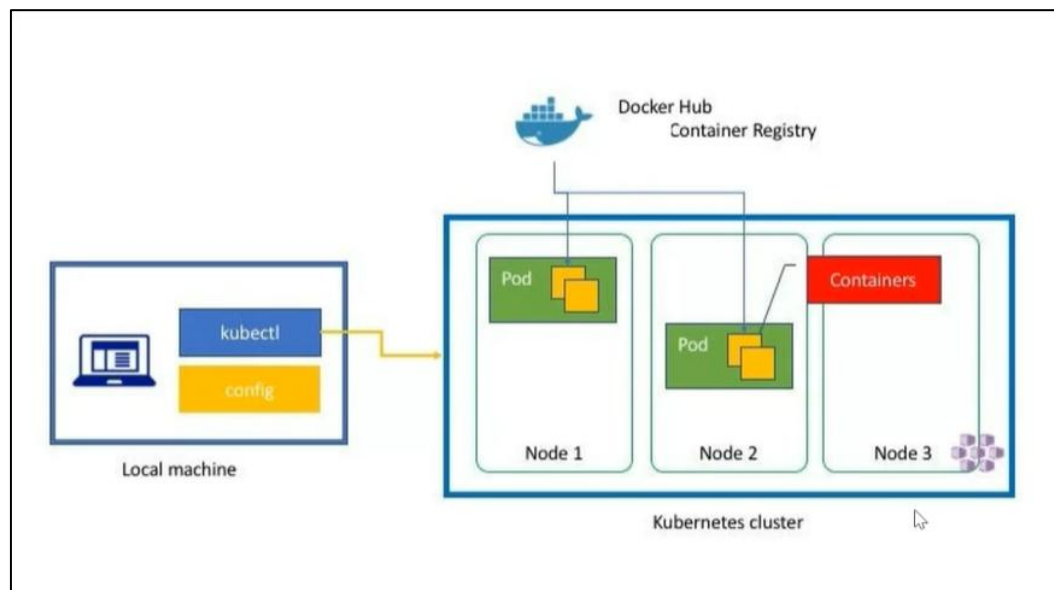
Ample Technologies is transforming its monolithic application into a microservices-based architecture. The application consists of one front-end service (Ruby) and two backend services (Node.js and Crystal). As an AWS Solutions Architect, I will deploy these services on AWS using Elastic Kubernetes Service (EKS), ensuring scalability and proper load balancing.

## 2. AWS Services Used:

- Amazon EKS
- Amazon EC2 (Node Groups)
- IAM (Roles and Policies)
- Elastic Load Balancer
- AWS CLI

## 3. Architecture Diagram

Below is the architecture diagram for the project:



## 4. Implementation Steps

## Step 0: Install AWS CLI

Install AWS CLI from [AWS CLI Installation Guide](#)

```
C:\Users\reesa>aws --version
aws-cli/2.23.2 Python/3.12.6 Windows/11 exe/AMD64
C:\Users\reesa>
```

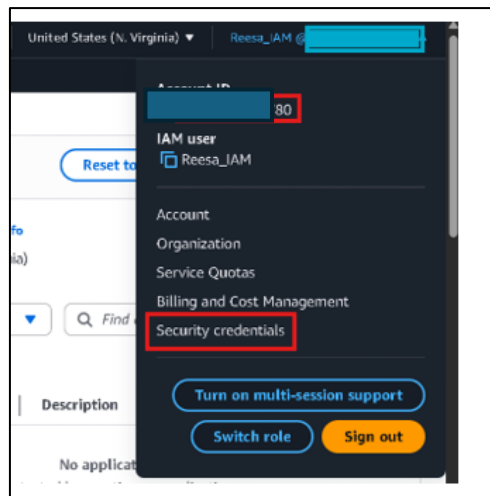
## Step 1: Configure AWS CLI & Create EKS Cluster

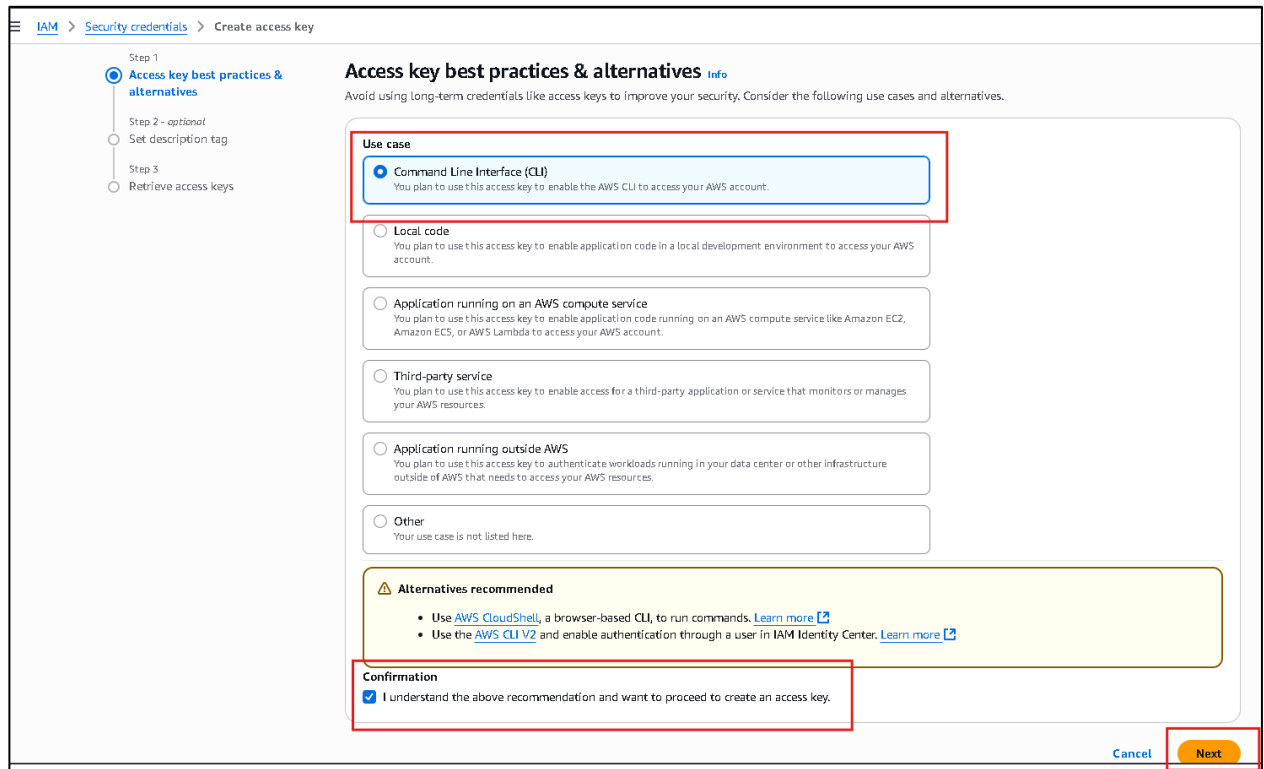
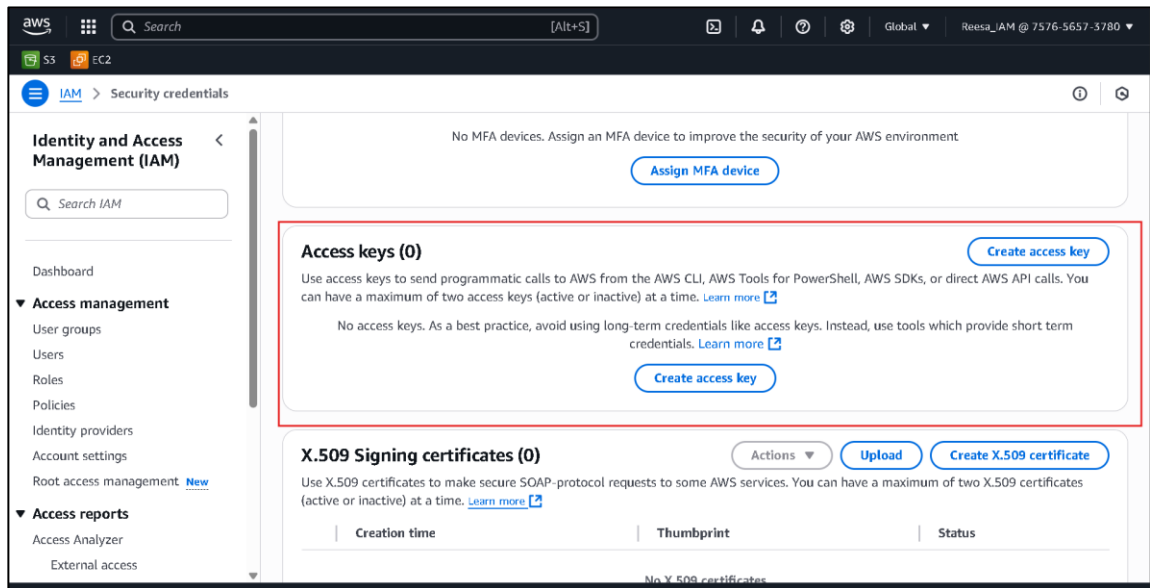
- a. Run *aws configure*

*aws configure*

Before running , retrieve your **AWS Access Key ID** and **Secret Access Key** by following these steps:

1. Go to AWS Console → IAM → Users
2. Click on your username
3. Navigate to the “**Security credentials**” tab
4. Click “**Create access key**”
5. Choose “**Command Line Interface (CLI)**” use case and click Next
6. Download the .csv file or copy the keys displayed





### Set description tag - optional [Info](#)

The description for this access key will be attached to this user as a tag and shown alongside the access key.

#### Description tag value

Describe the purpose of this access key and where it will be used. A good description will help you rotate this access key confidently later.

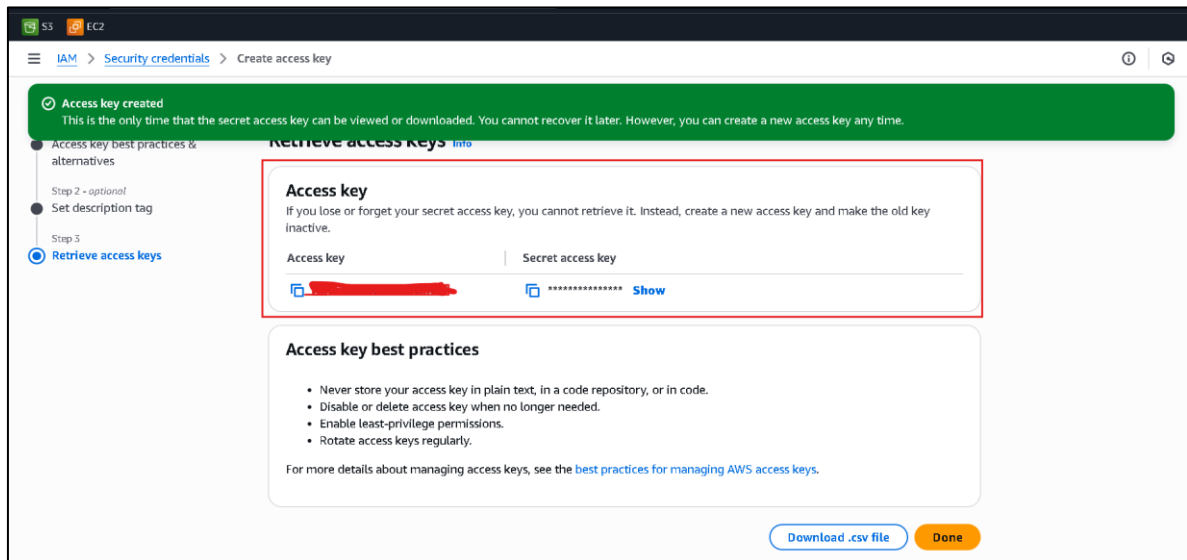
For Microservice project

Maximum 256 characters. Allowed characters are letters, numbers, spaces representable in UTF-8, and: \_ . : / = + - @

[Cancel](#)

[Previous](#)

[Create access key](#)



Then, open your terminal and run:

```
aws configure
```

Enter the following:

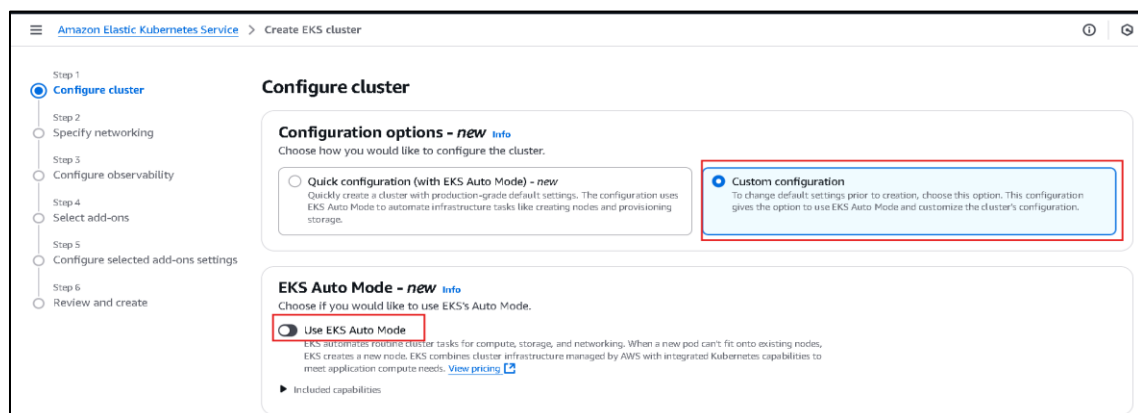
- **AWS Access Key ID**
- **AWS Secret Access Key**
- **Region** (e.g., us-east-2)
- **Output format** (e.g., json)

```
C:\Users\reesa>aws configure
AWS Access Key ID [*****7SSB]: 
AWS Secret Access Key [*****TFcr]: 
Default region name [us-east-2]: us-east-2
Default output format [json]: json

C:\Users\reesa>
```

## Step 2: Create EKS Cluster

1. Navigate to the EKS console.
2. Click on “**Add cluster**” and select “**Create**”.
3. Provide a name for your cluster (e.g., demoeks).



#### 4. Select and create the IAM role as shown below.



**Cluster configuration** [Info](#)

**Name**  
Enter a unique name for this cluster. This property cannot be changed after the cluster is created.

demoeks

The cluster name should begin with letter or digit and can have any of the following characters: the set of Unicode letters, digits, hyphens and underscores. Maximum length of 100.

**Cluster IAM role** [Info](#)  
Select the Cluster IAM role to allow the Kubernetes control plane to manage AWS resources on your behalf. This cannot be changed after the cluster is created. To create a new custom role, follow the instructions in the [Amazon EKS User Guide](#).

Select role ▼  **Create recommended role** 


**Identity and Access Management (IAM)**


Search IAM

Dashboard

▼ Access management

- User groups
- Users
- Roles**
- Policies
- Identity providers
- Account settings
- Root access management [New](#)

**Role eks-cluster created.** [View role](#) 

**Roles (11)** [Info](#)  [Delete](#) [Create role](#)

An IAM role is an identity you can create that has specific permissions with credentials that are valid for short durations. Roles can be assumed by entities that you trust.

Search

<input type="checkbox"/>	Role name	Trusted entities	Last activity
<input type="checkbox"/>	<a href="#">AWSServiceRoleForAPIGateway</a>	AWS Service: ops.apigateway (Service-Linkage)	-
<input type="checkbox"/>	<a href="#">AWSServiceRoleForSMSVoice</a>	AWS Service: sms-voice (Service-Linkage)	-
<input type="checkbox"/>	<a href="#">AWSServiceRoleForSupport</a>	AWS Service: support (Service-Linkage)	-
<input type="checkbox"/>	<a href="#">AWSServiceRoleForTrustedAdvisor</a>	AWS Service: trustedadvisor (Service-Linkage)	-
<input type="checkbox"/>	<b>eks-cluster</b>	AWS Service: eks	-



**Cluster configuration** [Info](#)

**Name**  
Enter a unique name for this cluster. This property cannot be changed after the cluster is created.

demoeks

The cluster name should begin with letter or digit and can have any of the following characters: the set of Unicode letters, digits, hyphens and underscores. Maximum length of 100.

**Cluster IAM role** [Info](#)  
Select the Cluster IAM role to allow the Kubernetes control plane to manage AWS resources on your behalf. This cannot be changed after the cluster is created. To create a new custom role, follow the instructions in the [Amazon EKS User Guide](#).

eks-cluster ▼  **Create recommended role** 

#### 5. Configure networking settings as required.


**Amazon Elastic Kubernetes Service** > Create EKS cluster

Step 1: Configure cluster  
Step 2: **Specify networking**  
Step 3: Configure observability  
Step 4: Select add-ons  
Step 5: Configure selected add-ons settings  
Step 6: Review and create


**Specify networking**

**Networking** [Info](#)  
IP address family and service IP address range cannot be changed after cluster creation.

**VPC** [Info](#)  
Select a VPC to use for your EKS cluster resources.

vpc-050fdd9ae67662f2f | Default 

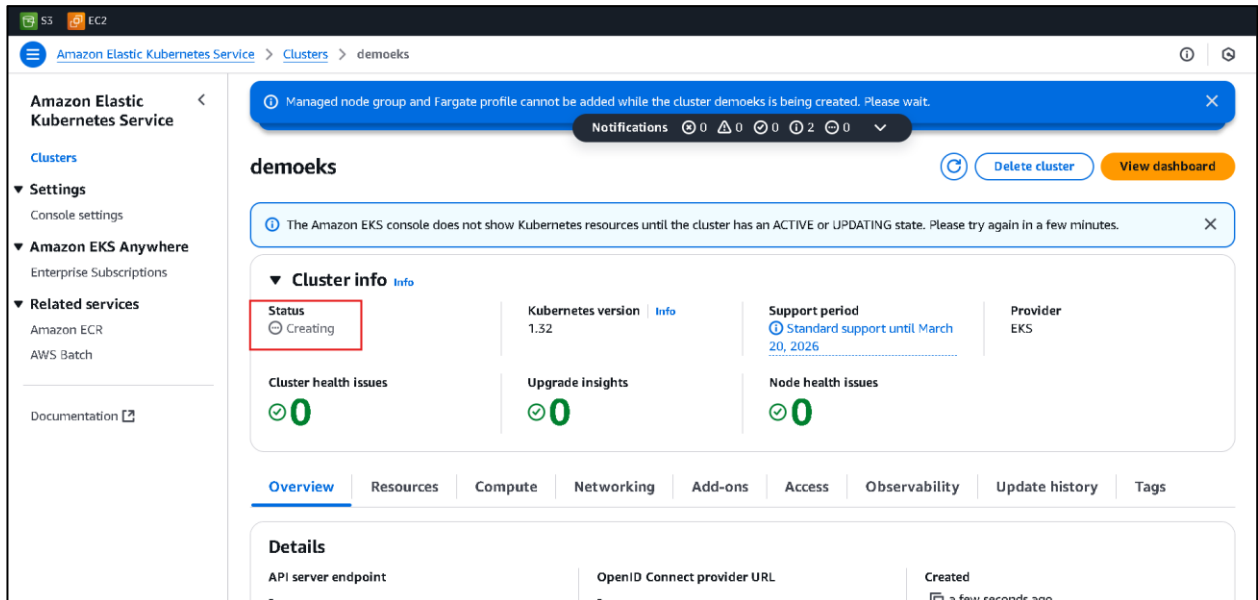
**Subnets** [Info](#)  
Choose the subnets in your VPC where the control plane may place elastic network interfaces (ENIs) to facilitate communication with your cluster. To create a new subnet, go to the corresponding page in the [VPC console](#).

Select subnets ▼  [Clear selected subnets](#)

subnet-0a3606af8f1e96852 us-east-1a 172.31.80.0/20	subnet-02b308584051da3cb us-east-1d 172.31.0.0/20
subnet-0b1de0b465b2943d9 us-east-1c 172.31.32.0/20	subnet-0365c10e4979aa754 us-east-1b 172.31.16.0/20
subnet-025c4b46deb96ecd2 us-east-1f 172.31.64.0/20	

The subnet in us-east-1f was removed during EKS cluster creation due to potential availability zone capacity issues or compatibility limitations. AWS occasionally restricts certain instance types or services in specific Availability Zones, so it's a best practice to use subnets in at least two validated and available AZs.

## 6. Create the cluster



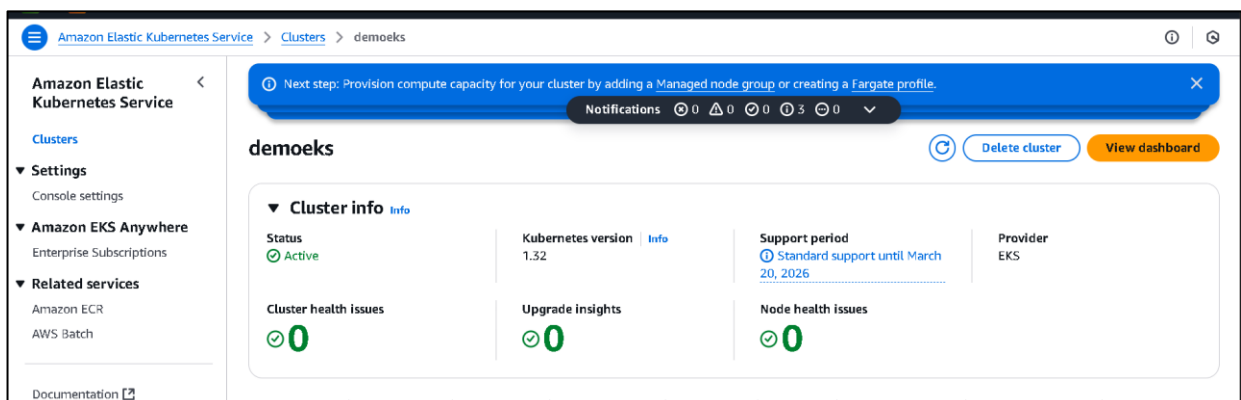
## Step 3: Update kubeconfig

- Run the following command to update your kubeconfig file with the new cluster's context:

```
aws eks --region us-east-1 update-kubeconfig --
```

```
C:\Users\reesa>aws eks --region us-east-1 update-kubeconfig --name demoeks
Cluster status is CREATING
C:\Users\reesa>
```

That means your EKS cluster wasn't fully created yet, so the kubeconfig couldn't be updated or used.



```
C:\Users\reesa>aws eks --region us-east-1 update-kubeconfig --name demoeks
Added new context arn:aws:eks:us-east-1:757656573780:cluster/demoeks to C:\Users\reesa\.kube\config
C:\Users\reesa>
```

This means the kubeconfig is now successfully set up and you can interact with your EKS cluster using kubectl.

#### Step 4: Install kubectl

- If kubectl is not installed, follow these steps:

```
curl.exe -LO "https://dl.k8s.io/release/v1.29.0/bin/windows/amd64/kubectl.exe"
curl.exe -LO "https://dl.k8s.io/v1.29.0/bin/windows/amd64/kubectl.exe.sha256"
CertUtil -hashfile kubectl.exe SHA256
type kubectl.exe.sha256
```

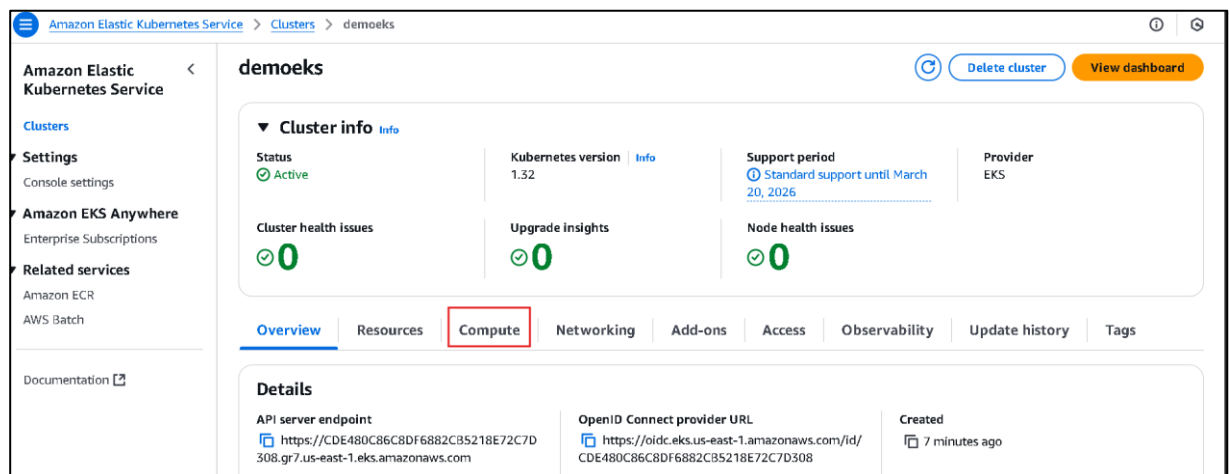
- Add the directory containing kubectl.exe to your system's PATH environment variable.
- Verify the installation:

```
kubectl version --client
```

```
C:\Users\reesa>kubectl version --client
Client Version: v1.31.4
Kustomize Version: v5.4.2
```

#### Step 5: Create Node Group

1. In the EKS console, navigate to your cluster.
2. Click on “**Add node group**”.



3. Provide a name for the node group.
4. Select and create an IAM role as shown below.

5. Refresh and add the created IAM role

6. Configure the following settings:

- Instance type: t2.micro
- Disk size: 5 GB
- Minimum size: 3
- Maximum size: 3
- Desired size: 3



## Node group scaling configuration

### Desired size

Set the desired number of nodes that the group should launch with initially.

3 nodes

Desired node size must be greater than or equal to 0

### Minimum size

Set the minimum number of nodes that the group can scale in to.

3 nodes

Minimum node size must be greater than or equal to 0

### Maximum size

Set the maximum number of nodes that the group can scale out to.

3 nodes

Maximum node size must be greater than or equal to 1 and cannot be lower than the minimum size

7. Allow remote access and select an existing SSH key pair.

8. Create the node group.

The screenshot shows the Amazon Elastic Kubernetes Service (EKS) console. The left sidebar contains navigation links for Amazon Elastic Kubernetes Service, Clusters, Settings, Amazon EKS Anywhere, and Related services. The main content area shows the details for the 'demoeks' cluster. The 'Compute' tab is selected, displaying a table of 3 nodes. The table has columns for Node name, Instance type, Compute, Managed by, Created, and Status. All three nodes are in 'Ready' status.

Node name	Instance type	Compute	Managed by	Created	Status
ip-172-31-1-37.ec2.internal	t3.medium	Node group	demo-eks-nodes	a few seconds ago	Ready
ip-172-31-32-56.ec2.internal	t3.medium	Node group	demo-eks-nodes	a few seconds ago	Ready
ip-172-31-85-59.ec2.internal	t3.medium	Node group	demo-eks-nodes	a few seconds ago	Ready

## Step 6: Verify Node Group

➤ Run the following command to verify that the nodes are active:

```
kubectl get nodes
```

```
C:\Users\reesa>kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-1-37.ec2.internal         Ready    <none>    109s  v1.32.3-eks-473151a
ip-172-31-32-56.ec2.internal        Ready    <none>    111s  v1.32.3-eks-473151a
ip-172-31-85-59.ec2.internal        Ready    <none>    110s  v1.32.3-eks-473151a

C:\Users\reesa>
```

## Step 7: Deploy Microservices

This application consists of three microservices, each containerized and deployed independently on Amazon EKS. The Kubernetes manifests for these services are provided in the zipped folder available on my GitHub repository inside the 'Code' file (ecsdemo-nodejs, ecsdemo-crystal, and ecsdemo-frontend).

### 1. Deploy ecsdemo-nodejs Service

- A Node.js-based backend API.
- Exposes an internal service (ClusterIP) for communication within the Kubernetes cluster.
- Does **not** require public access or a Load Balancer.
  - Navigate to the ecsdemo-nodejs/kubernetes directory.
  - Apply the deployment and service manifests:

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-nodejs\kubernetes>kubectl apply -f deployment.yaml
deployment.apps/ecsdemo-nodejs created

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-nodejs\kubernetes>kubectl apply -f service.yaml
service/ecsdemo-nodejs created

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-nodejs\kubernetes>
```

### 2. Deploy ecsdemo-crystal Service

- A backend API written in Crystal programming language.
- Also exposed internally via ClusterIP.
- Used by the frontend to fetch or process data.

1. Navigate to the ecsdemo-crystal/kubernetes directory.
2. Apply the deployment and service manifests:

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes>kubectl apply -f deployment.yaml
deployment.apps/ecsdemo-crystal created

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes>kubectl apply -f service.yaml
service/ecsdemo-crystal created

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes>
```

### 3. Deploy ecsdemo-frontend Service

- The user-facing component built with Ruby.
- Exposed publicly via a Kubernetes LoadBalancer service.
- Automatically provisions an **Elastic Load Balancer (ELB)**, allowing access to the application from the internet using the provided DNS name.

1. Navigate to the ecsdemo-frontend/kubernetes directory.
2. Apply the deployment and service manifests:

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-frontend\kubernetes>kubectl apply -f deployment.yaml
deployment.apps/ecsdemo-frontend created

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-frontend\kubernetes>kubectl apply -f service.yaml
service/ecsdemo-frontend created

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-frontend\kubernetes>
```

### 3. Verify the deployment and get the service details

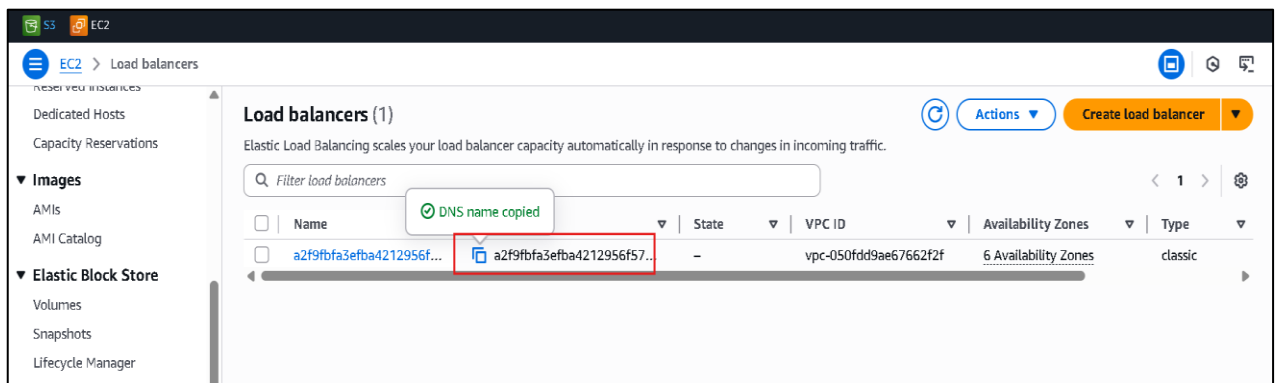
```
kubectl get pods
```

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes>kub
ectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
ecsdemo-crystal-76796bbc-f-2rnrl    1/1     Running   0           43s
ecsdemo-frontend-54d56d89fb-76sp9   1/1     Running   0           7m19s
ecsdemo-nodejs-6d679857c4-9cbpn     1/1     Running   0           2m19s
```

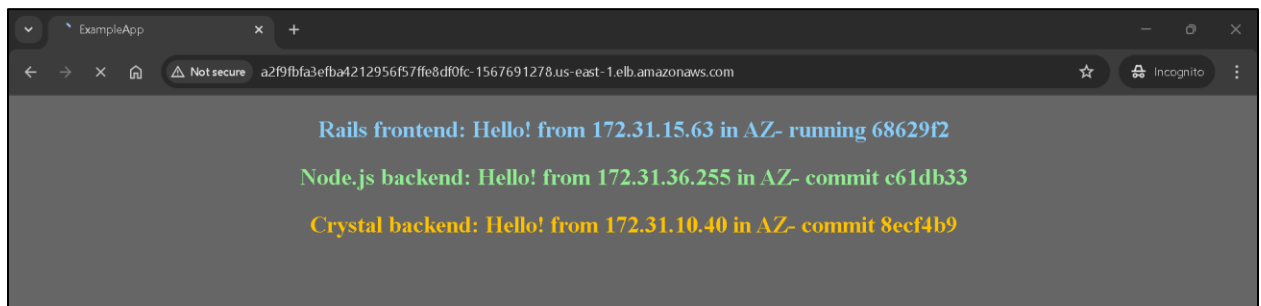
## 5. Testing & Validation

### Step 1: Open the Application via Load Balancer

1. Go to the **EC2 Console** → **Load Balancers** (under **Load Balancing** in the left sidebar).
2. Click on the load balancer associated with your **ecsdemo-frontend** service.
3. Copy the **DNS name** of the load balancer.



4. Open an **Incognito window** in your browser and **paste the DNS name** in the address bar.



5. You should see the deployed application running.

### Step 2: Scale Deployments

- Scale each deployment to 3 replicas:

```
kubectl scale deployment ecsdemo-nodejs --replicas=3
kubectl scale deployment ecsdemo-crystal --replicas=3
kubectl scale deployment ecsdemo-frontend --replicas=3
```

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes> kubectl scale deployment.ecsdemo-crystal --replicas=3
deployment.apps/ecsdemo-crystal scaled

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes> get pods
'get' is not recognized as an internal or external command,
operable program or batch file.

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
ecsdemo-crystal-76796bbcf-2nrnl	1/1	Running	0	4m34s
ecsdemo-crystal-76796bbcf-4t528	1/1	Running	0	18s
ecsdemo-crystal-76796bbcf-zvbvr	1/1	Running	0	18s
ecsdemo-frontend-54d56d89fb-76sp9	1/1	Running	0	11m
ecsdemo-nodejs-6d679857c4-9cbpn	1/1	Running	0	6m10s

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes>
```

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-crystal\kubernetes> cd C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-frontend\kubernetes
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-frontend\kubernetes> kubectl scale deployment.ecsdemo-frontend --replicas=3
deployment.apps/ecsdemo-frontend scaled

C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-frontend\kubernetes> cd C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-nodejs\kubernetes
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-nodejs\kubernetes> kubectl scale deployment.ecsdemo-nodejs --replicas=3
deployment.apps/ecsdemo-nodejs scaled
```

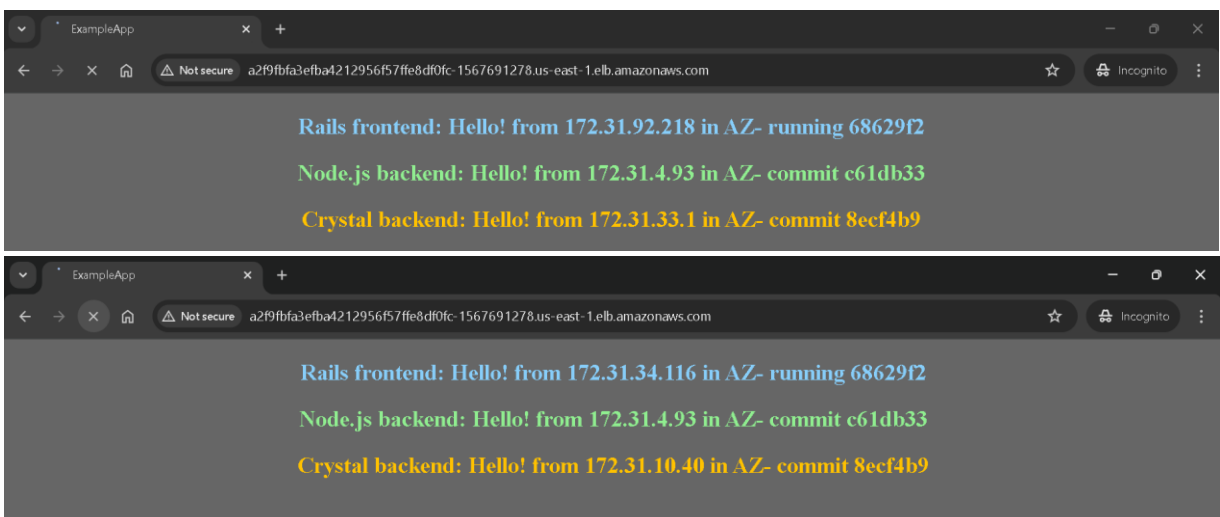
- After scaling the microservices using kubectl scale, Kubernetes automatically schedules additional pod replicas across the available worker nodes.

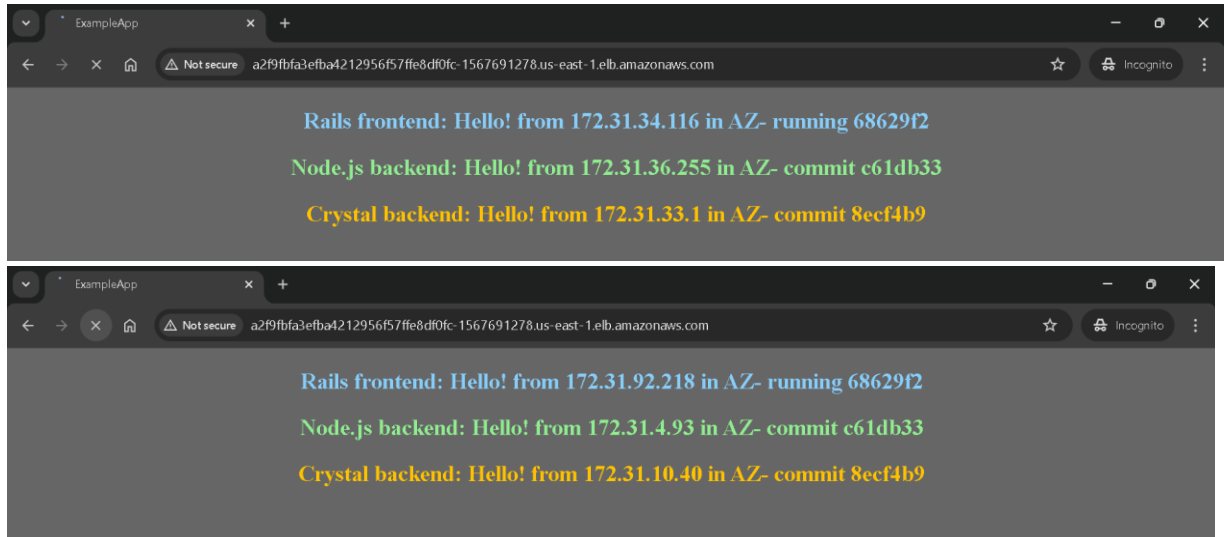
```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-nodejs\kubernetes> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
ecsdemo-crystal-76796bbcf-2nrnl	1/1	Running	0	7m54s
ecsdemo-crystal-76796bbcf-4t528	1/1	Running	0	3m38s
ecsdemo-crystal-76796bbcf-zvbvr	1/1	Running	0	3m38s
ecsdemo-frontend-54d56d89fb-76sp9	1/1	Running	0	14m
ecsdemo-frontend-54d56d89fb-kbkd5	1/1	Running	0	102s
ecsdemo-frontend-54d56d89fb-rsb7q	1/1	Running	0	102s
ecsdemo-nodejs-6d679857c4-9cbpn	1/1	Running	0	9m30s
ecsdemo-nodejs-6d679857c4-v8msk	1/1	Running	0	66s
ecsdemo-nodejs-6d679857c4-x2sqm	1/1	Running	0	66s

```
C:\Users\reesa\OneDrive\Desktop\AWS WORKS\DevOps Interview Questions\Kubernetes\6. Microservices\Codes\ecsdemo-nodejs\kubernetes>
```

- As a result, running will display multiple pods for each microservice, each with a **unique IP address** assigned by the Kubernetes network. This ensures **high availability** and **load distribution** across replicas.





## 6. Clean Up AWS Resources

To avoid incurring unnecessary charges, delete the following resources after the project is completed:

- **EKS Cluster**
- **Node Group**
- **IAM Roles and Policies** created for EKS
- **Load Balancer** associated with the frontend service
- **Key pairs**
- **Security Groups**
- **Access Keys**

## 7. Conclusion

This project successfully demonstrates the deployment of a real-world microservices-based architecture on Amazon EKS using Docker and Kubernetes. By containerizing and orchestrating three services — frontend, Node.js backend, and Crystal backend — the application was made scalable, modular, and cloud-native. Key AWS services like IAM, EC2, Load Balancer, and EKS were integrated to ensure secure, reliable, and high-performing deployments.

From configuring AWS CLI and Kubernetes tools to managing infrastructure and scaling deployments, this project offers a hands-on understanding of cloud-native DevOps practices. It showcases not only technical proficiency in using AWS and Kubernetes but also an ability to architect scalable, production-ready solutions aligned with modern application development standards.

This end-to-end experience strengthens your foundation in cloud engineering and DevOps, making you better equipped for roles involving cloud infrastructure automation and microservices deployment on AWS.